

The rod of Asclepios

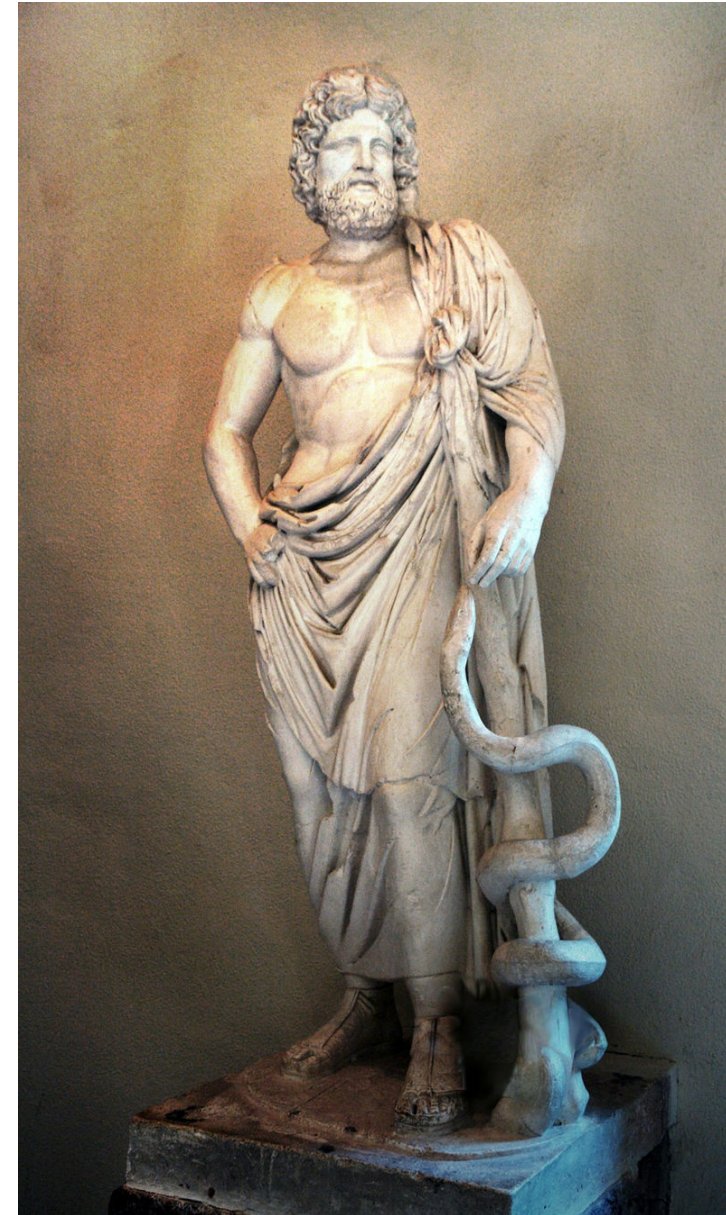
Jan Margeta | jan@kardio.me | [@jmargeta](https://twitter.com/jmargeta)

KardioMe³

**World full of
superpowers**

Python as our superpower?

Meet Asclepios



Hi, I am Jan

- Pythonista for 8 years
- Founder of KardioMe
- Maker of tools to better understand our hearts
- Python 3.5+, numpy, MxNet, Keras, Tensorflow, scikit-learn, SimpleITK, pydicom, Flask, Sanic, Django, PostgreSQL, ReactJS, Docker

This talk

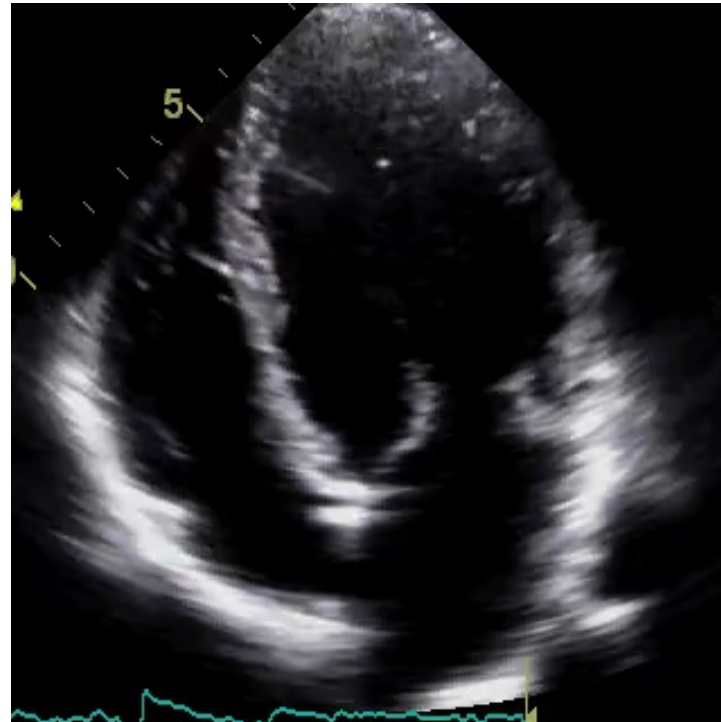
- Peek into our hearts with medical images
- What used to be hard is now magically simple
- Tricks and tools for machine learning in Python
we've learned along the way

Imaging of our hearts

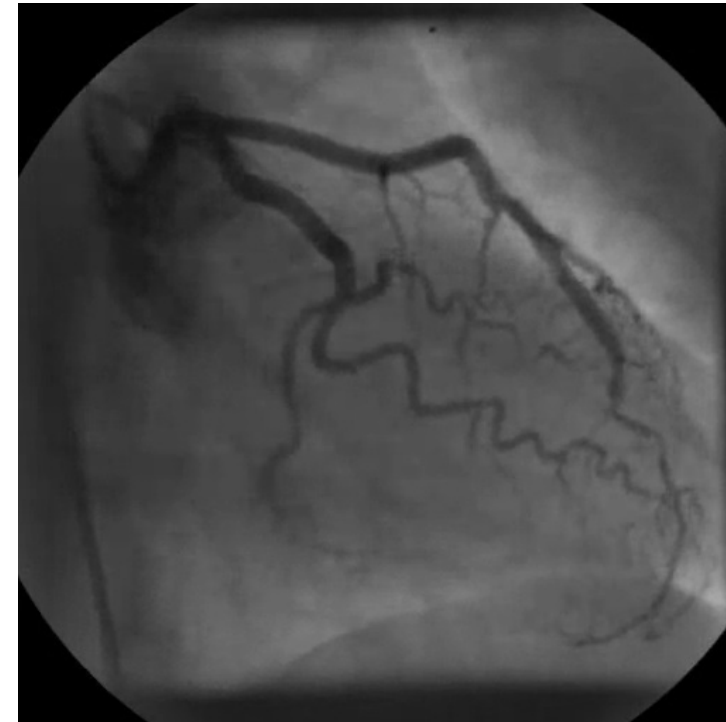
X-Ray



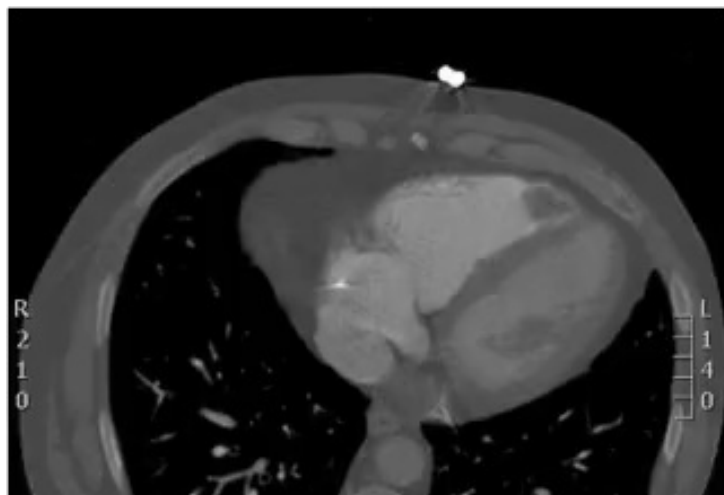
ultrasound



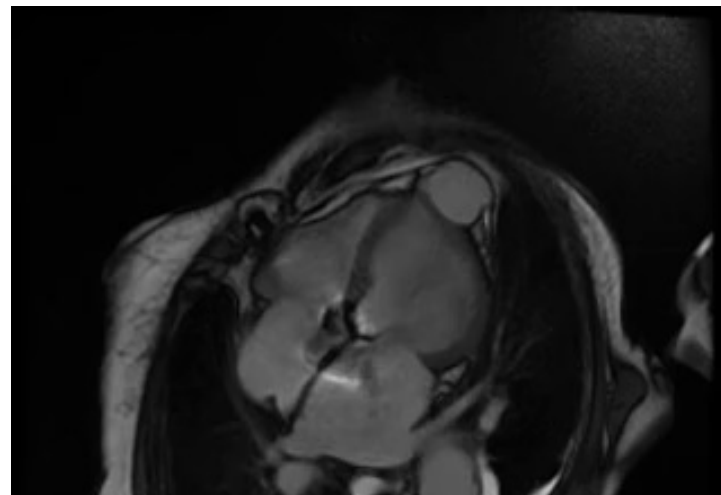
fluoroscopy



computed tomography



magnetic resonance



Kelly 2007
Carmo et al. 2010
Arnold et al. 2008
Foley et al. 2010
Vanezis et al. 2011



**What a human
sees**

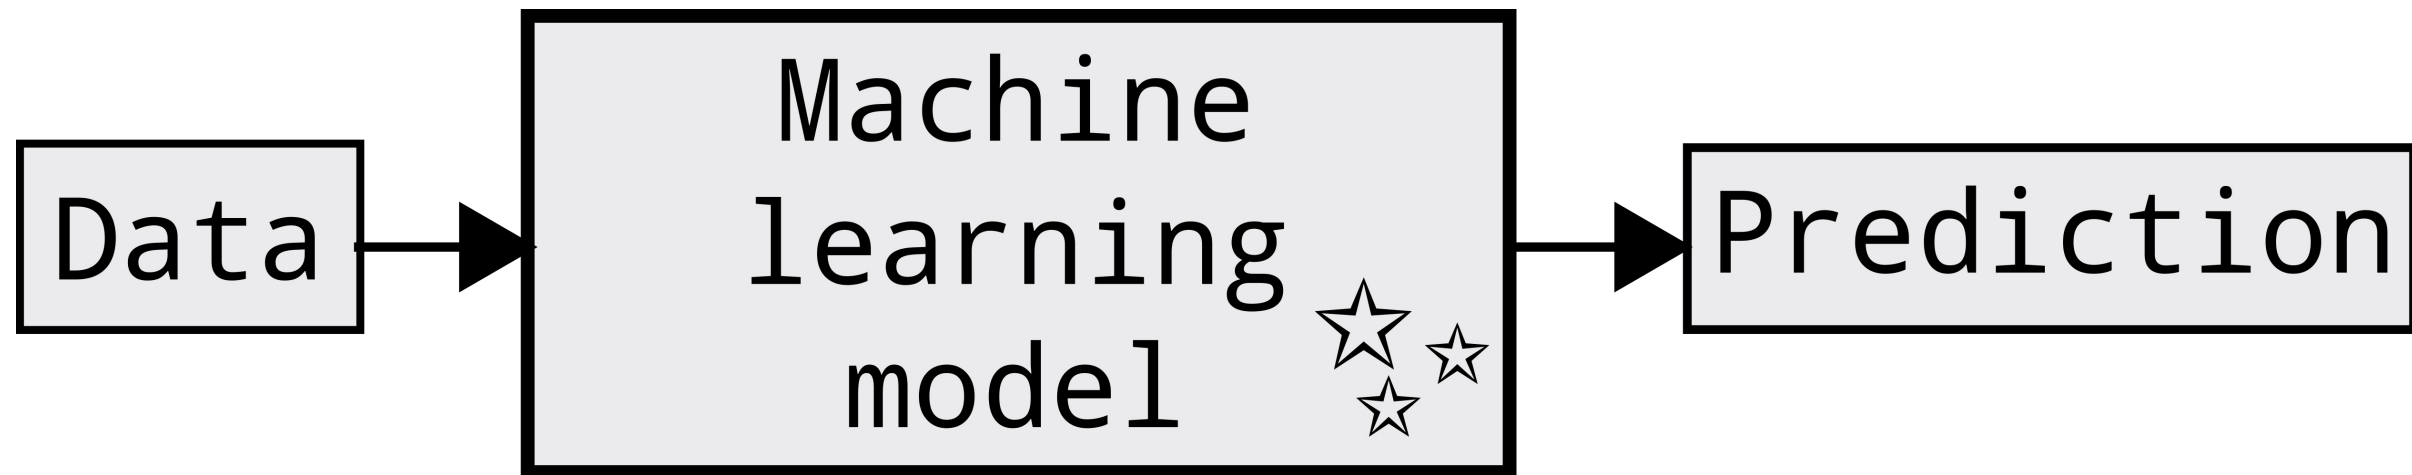


**What a human
sees**

133	115	57	65	52	69	129	145	134	121	106	81	87	59	42	44	18	24
131	87	53	57	140	143	133	130	119	117	112	115	81	94	48	31	31	25
139	59	54	120	138	135	137	139	120	110	107	109	122	68	105	33	36	16
119	60	46	144	138	138	131	133	120	111	67	45	46	98	91	44	37	18
106	47	71	151	121	115	127	122	116	107	77	88	88	61	77	43	32	17
103	43	80	125	127	124	120	109	120	109	106	110	56	83	92	35	34	24
90	51	59	118	123	113	107	111	119	101	99	91	108	103	69	43	31	32
101	60	59	105	105	105	98	99	105	109	95	114	101	110	54	39	36	22
90	45	49	94	101	96	94	109	104	100	107	93	98	39	49	34	31	18
67	46	45	58	95	102	77	85	101	79	91	97	71	58	43	38	16	19
19	39	50	48	54	72	98	91	68	88	97	79	61	36	33	19	17	20

Machine learning

Solving problems with data



prediction = model.*predict* (data)

Image recognition in 6 lines of code

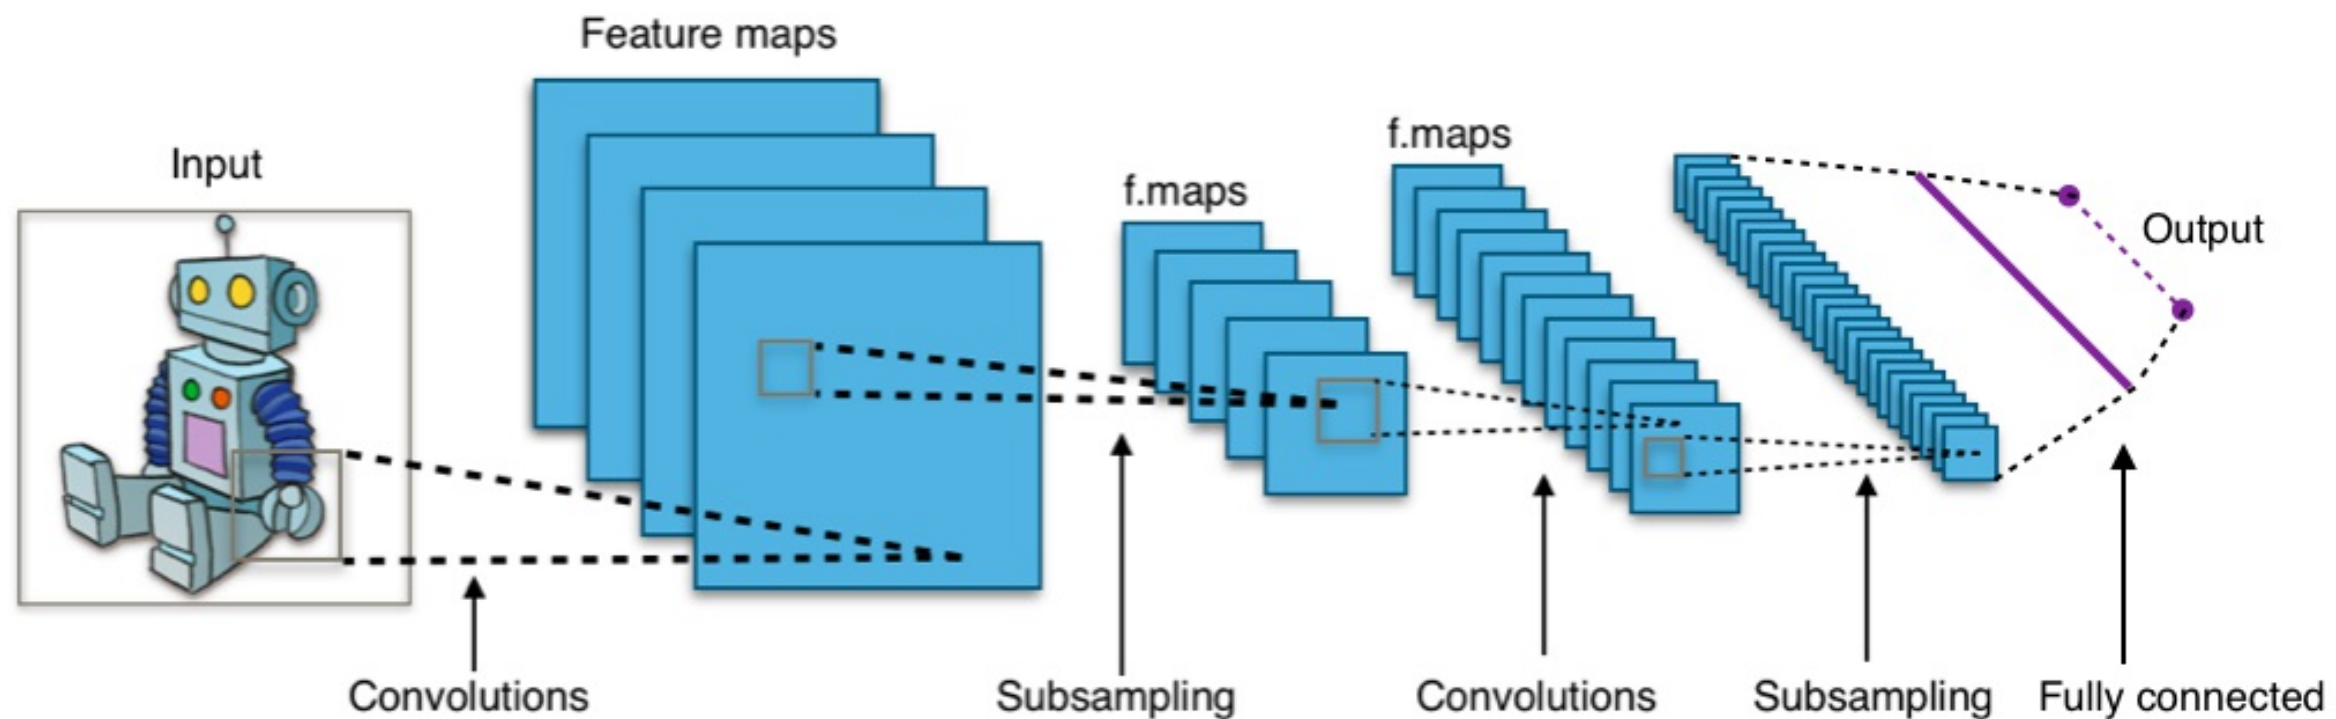
```
from keras.applications import imagenet_utils
from keras.applications.vgg16 import VGG16

# Load and prepare input images
images_raw = load_images()
images = imagenet_utils.preprocess_input(images_raw)

# Load a pretrained image classification model
model = VGG16(include_top=True, weights='imagenet')

# Do the prediction
predictions = model.predict(images)
```

Convolutional neural networks



By Aphex34

Finding the right image representation

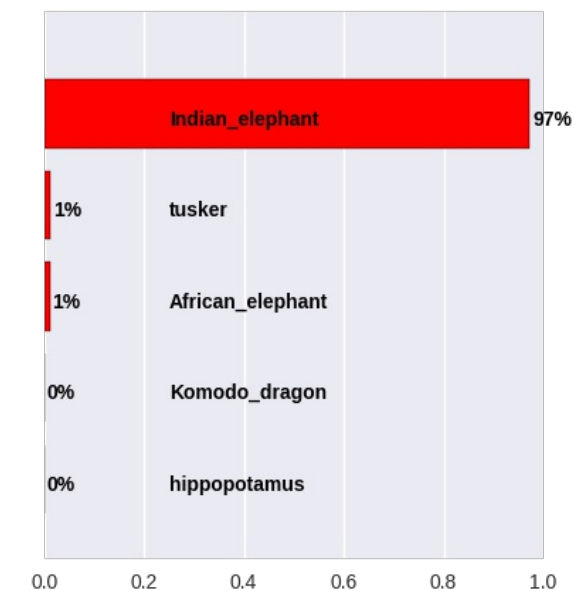
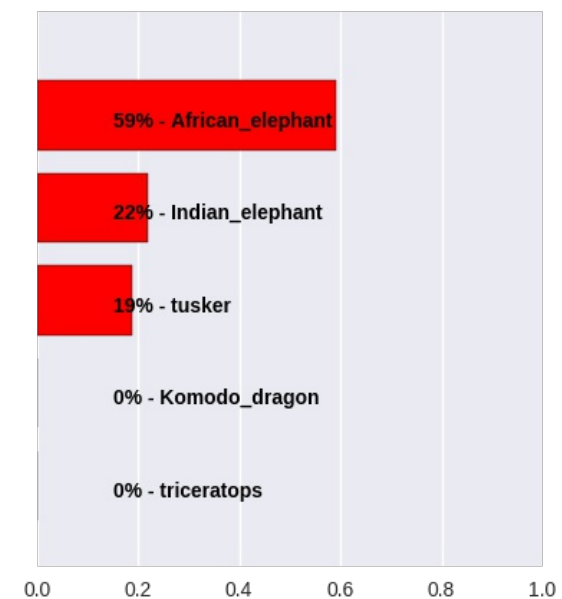
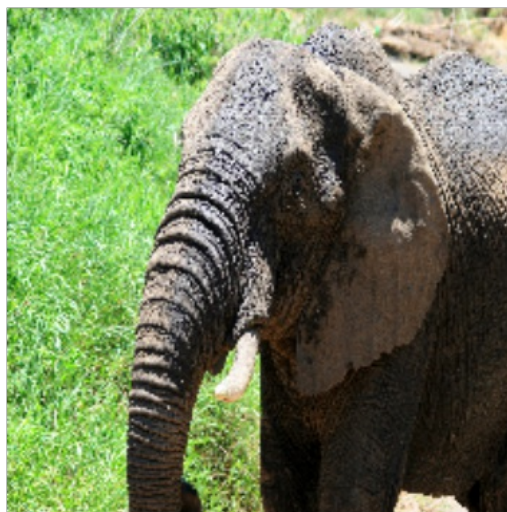
A black and white photograph of a squirrel, likely a red squirrel, shown in profile facing right. The squirrel's ears are large and pointed upwards, with a distinct white patch on the inner side of each ear. Its eyes are dark and focused forward. The fur appears soft and textured. The background is dark and out of focus, with some light-colored, blurry shapes that could be leaves or branches.



Excellent for natural images

Trained on Imagenet large scale visual recognition challenge dataset

10 million images, 1000 categories



Extracting visual features

```
# Load a pretrained classification model
source_model = VGG16(weights='imagenet')

# Define feature extractor from one layer of the network
feature_layer = source_model.get_layer('conv4')
feature_extractor = Model(
    input=fix_model.input,
    output=feature_layer.output)

# Extract features
features = feature_extractor.predict(images)
```

See also ["Deep visualization toolbox" on youtube](#)

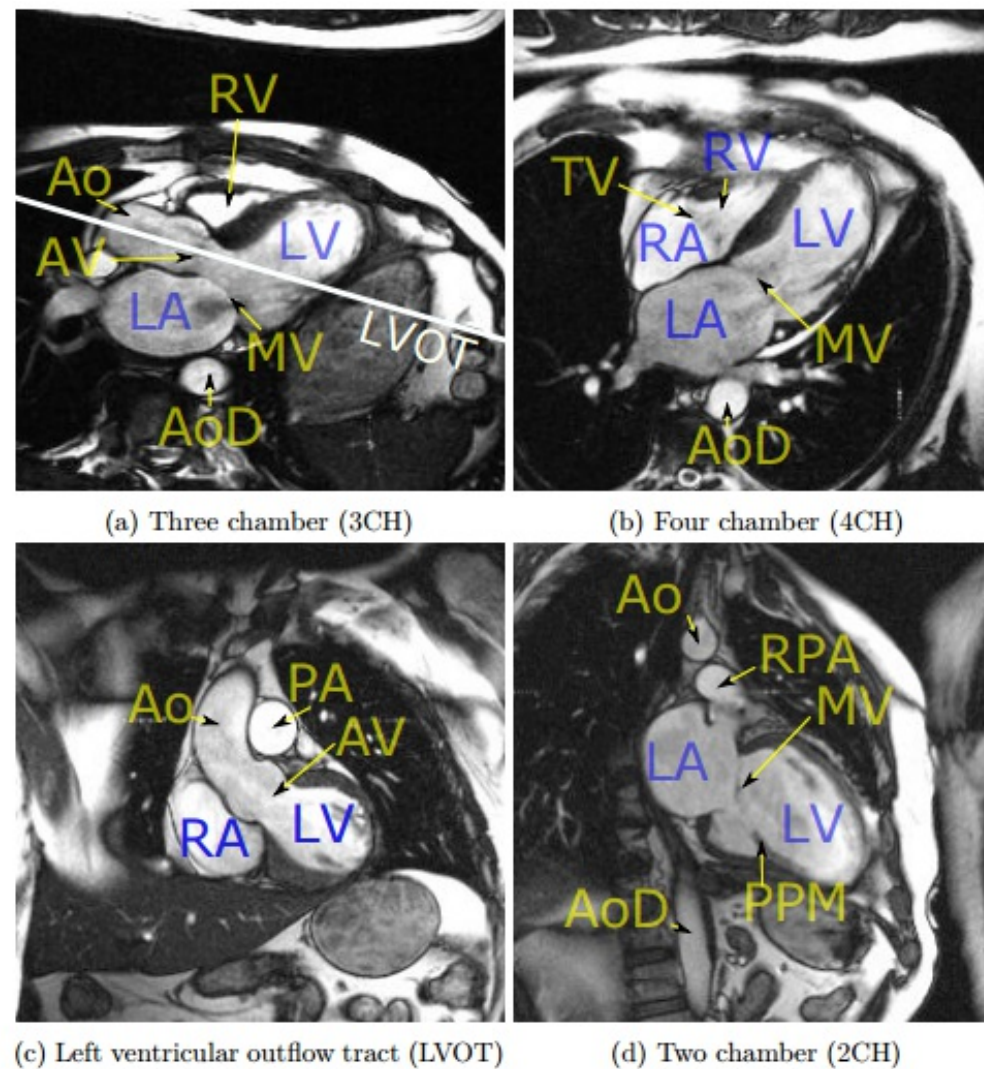
Using the extracted features with scikit-learn

```
from sklearn.svm import LinearSVC
def flatten_features(features):
    return features.reshape(len(features), -1)

features_train = feature_extractor.predict(images_train)
features_train = flatten_features(features_train)
classifier = LinearSVC()
classifier.fit(features_train, labels_train)

# predict on never seen images
features_test = feature_extractor.predict(images_test)
features_test = flatten_features(features_test)
prediction_test = classifier.predict(features_test)
```


Example: Cardiac view recognition



Train the model from scratch

```
from keras.models import Sequential
from keras.layers import Conv2D, Dense, Flatten

images_train, labels_train = load_data()
shape = (64, 64, 1)

model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=shape),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(4, activation='softmax'),
])
```

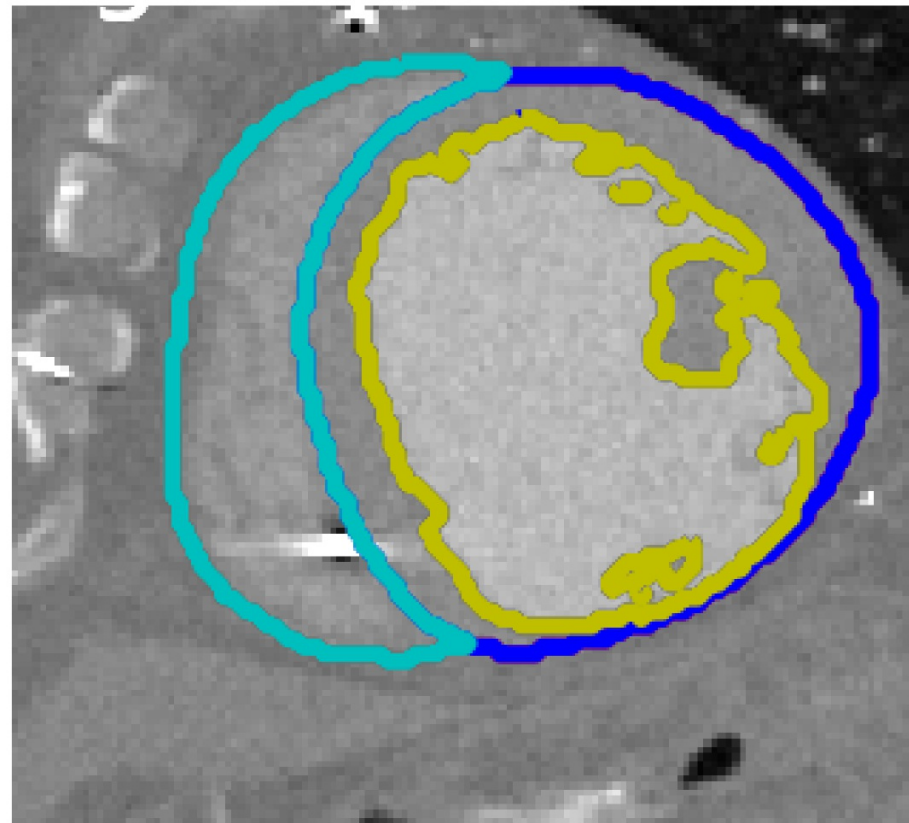
```
# Loss function - task dependent
# high for bad parameters, low for good ones
# e.g. for image recognition
loss_function = 'sparse_categorical_crossentropy'

# Compile the model and fit
model.compile(loss=loss_function, optimizer='adam')
model.fit(images_train, labels_train)

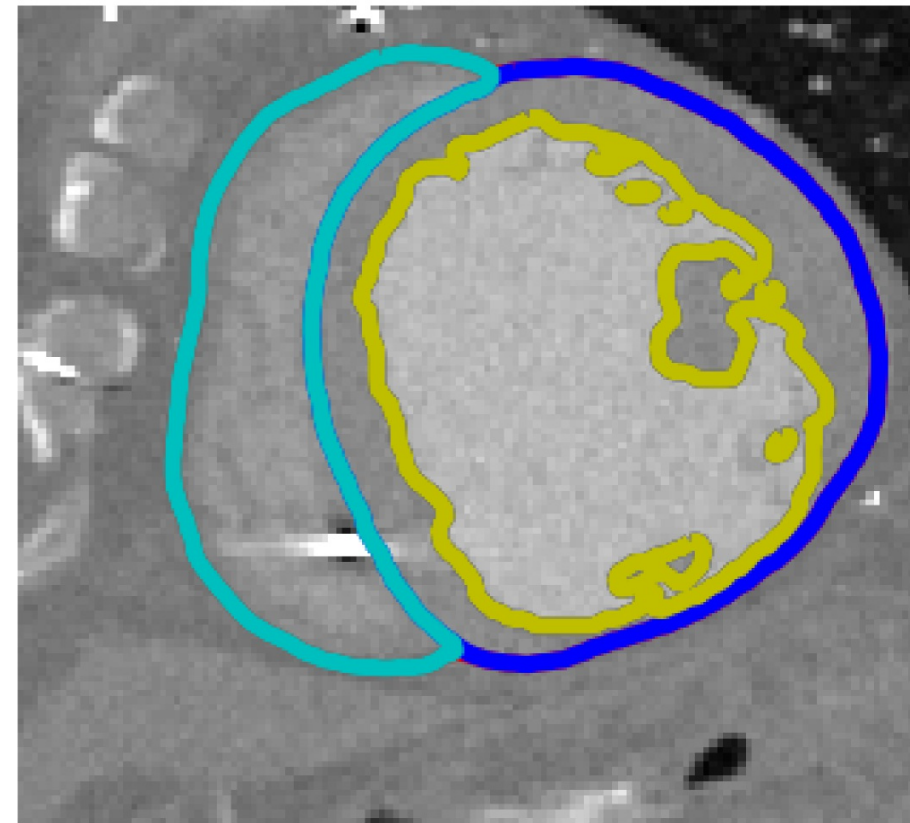
# Save the model for reuse
model.save('model.h5')
```

Let's save some time for our radiologists

30 minutes



12 seconds

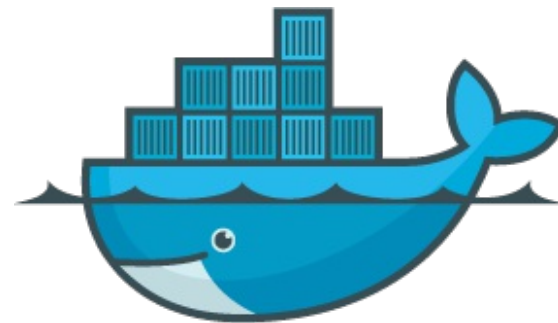


Deploy as a web service



Flask

web development,
one drop at a time



docker

see also Tensorflow-serving and Kubernetes

Expose the model with Flask*

```
import keras
from flask import Flask, jsonify, request
app = Flask(__name__)
model = keras.models.load_model('model.h5')

@app.route('/predict', methods=['POST'])
def predict():
    image_batch = request_to_numpy(request)
    y = model.predict(image_batch)
    prediction = convert_prediction(y)
    return jsonify(output=prediction)

app.run(port=5000, threaded=False)
```

*do not run in production, it requires a bit more love than this

**Run with the same conditions as
when it was built**



Define the Dockerfile

```
FROM python:3.5

RUN mkdir -p /usr/src/app
COPY server.py /usr/src/app/
COPY model.h5 /usr/src/app/
COPY requirements.txt /usr/src/app/
WORKDIR /usr/src/app
RUN pip install -r requirements.txt

EXPOSE 5000
CMD python server.py
```


Build the Docker container

```
docker build -t kardiome/model-pyparis .
```

Run the service

```
docker run -d -p 5000:5000 kardiome/model-pycon
```

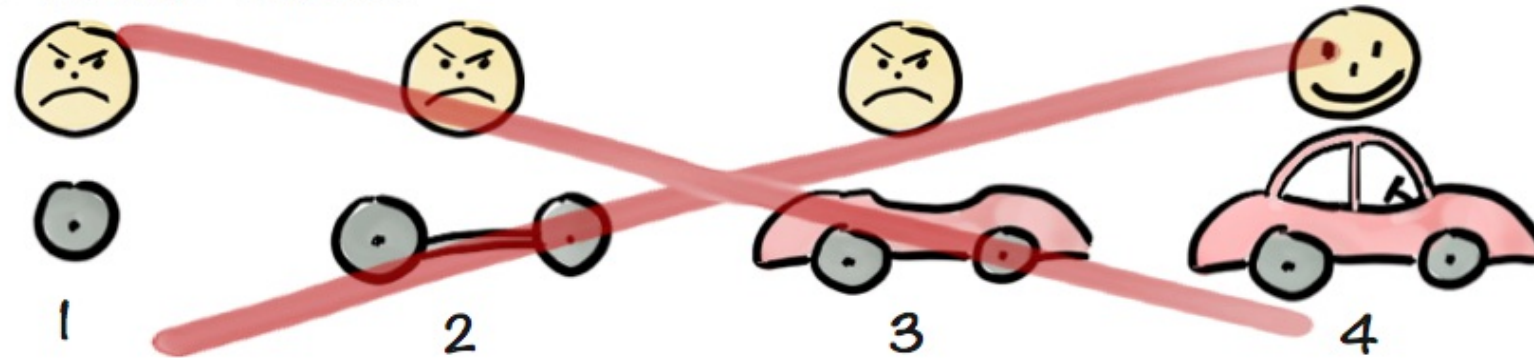
Call the service

```
curl -X POST -F 'image=@/data/im.png' localhost:5000/predict
```

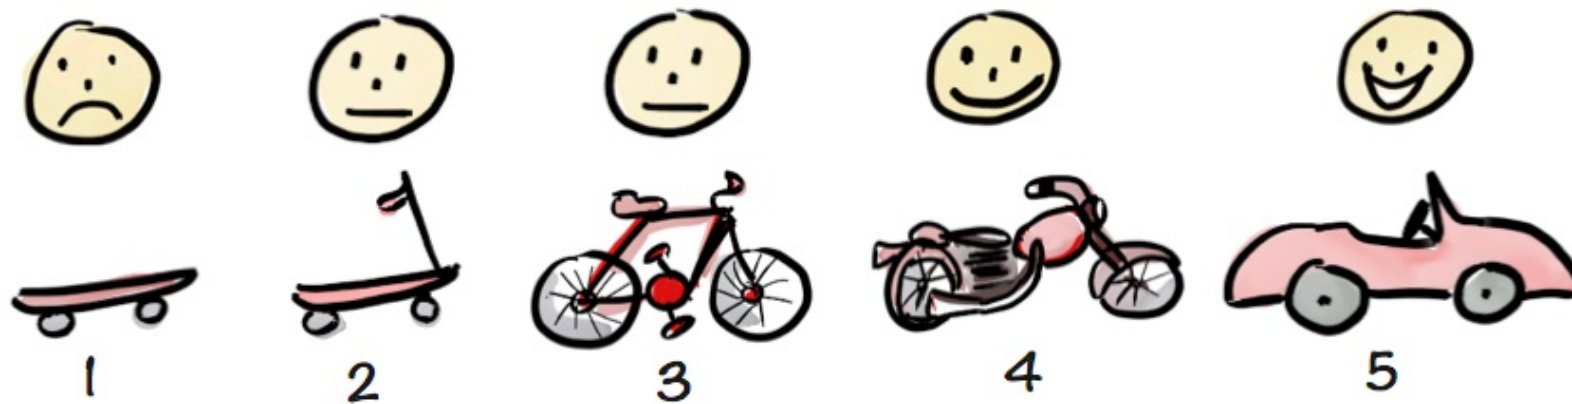
**Tips to improve
your machine
learning today**

Iterate fast

Not like this....



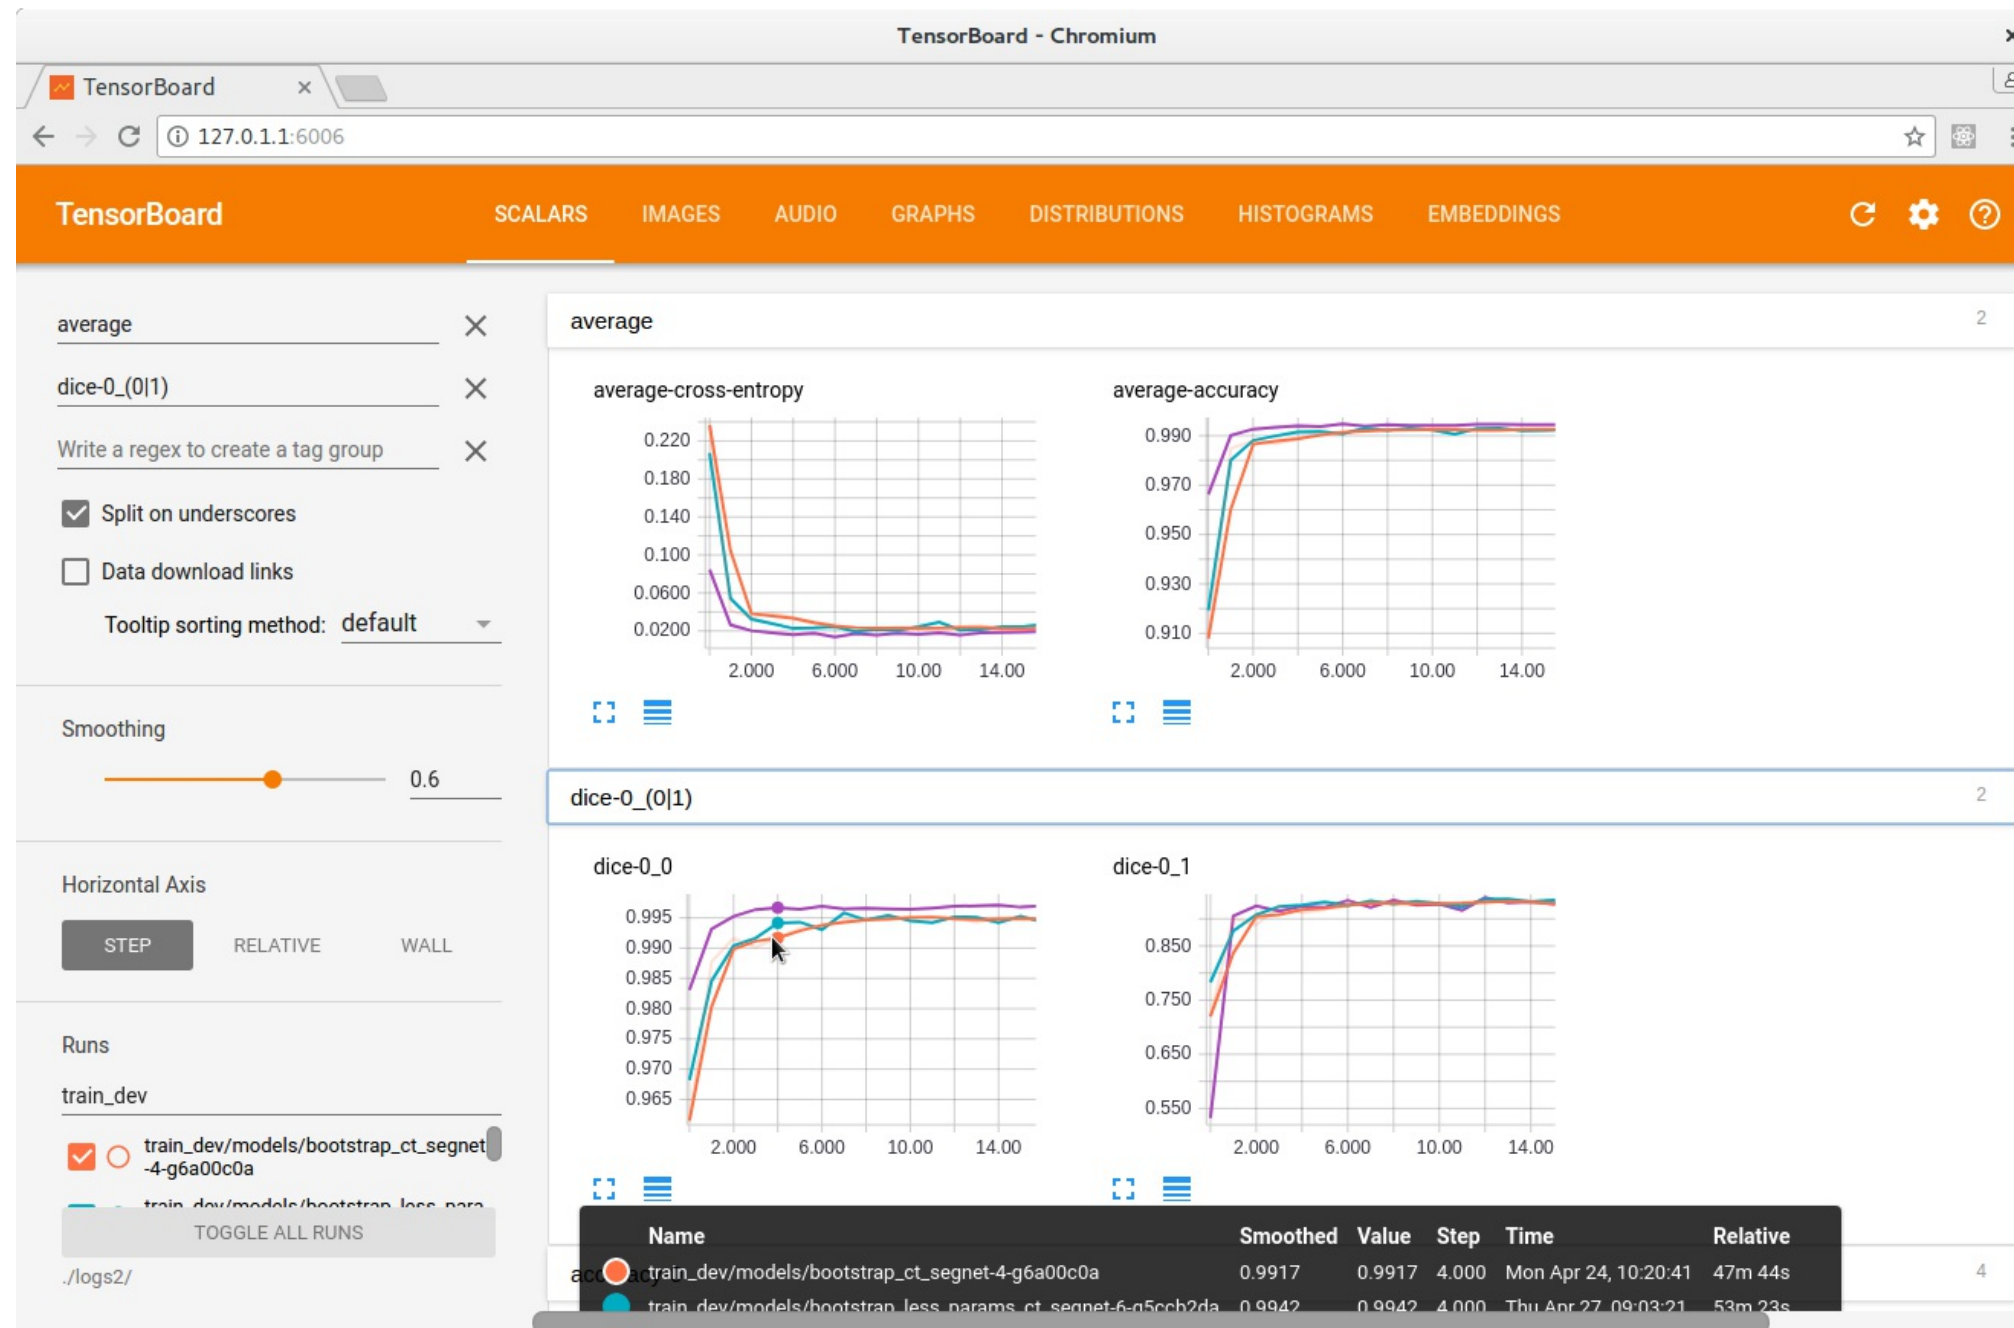
Like this!



by Henrik Kniberg

Minimally viable model

One metric to rule them all



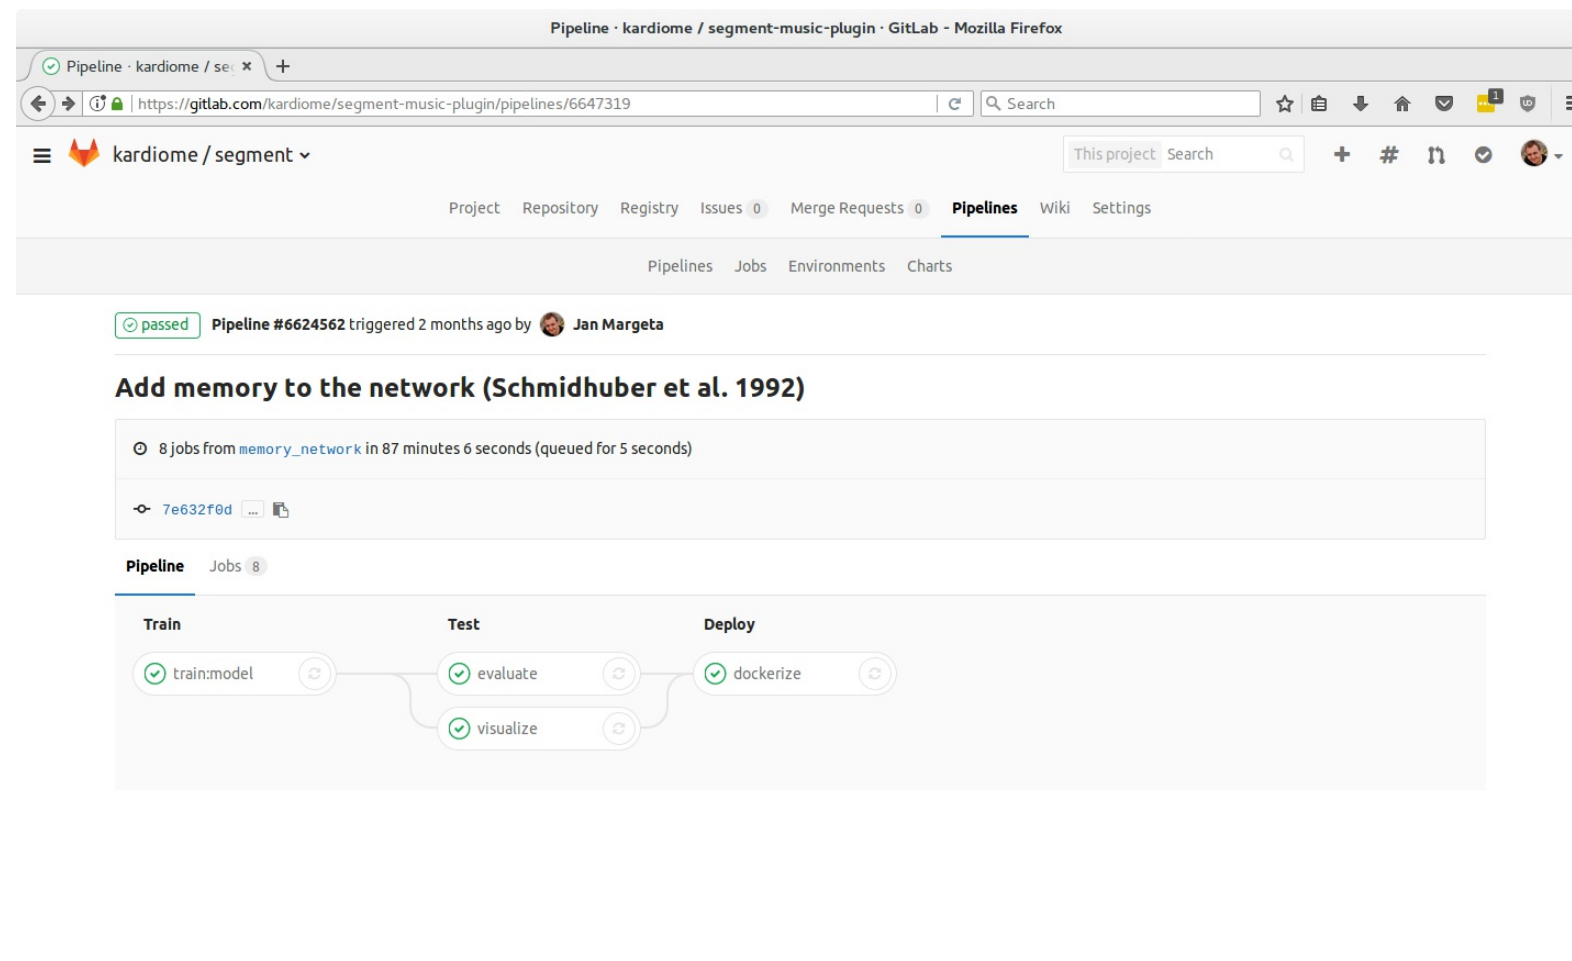
Training progress visualisation in Tensorboard

Visualize everything



Built with Sanic backend (like Flask) and React

Progress with confidence and repeatable pipelines



Gitlab's continuous integration is an excellent start for simple pipelines
(see Airflow, Luigi, Joblib, Flink)

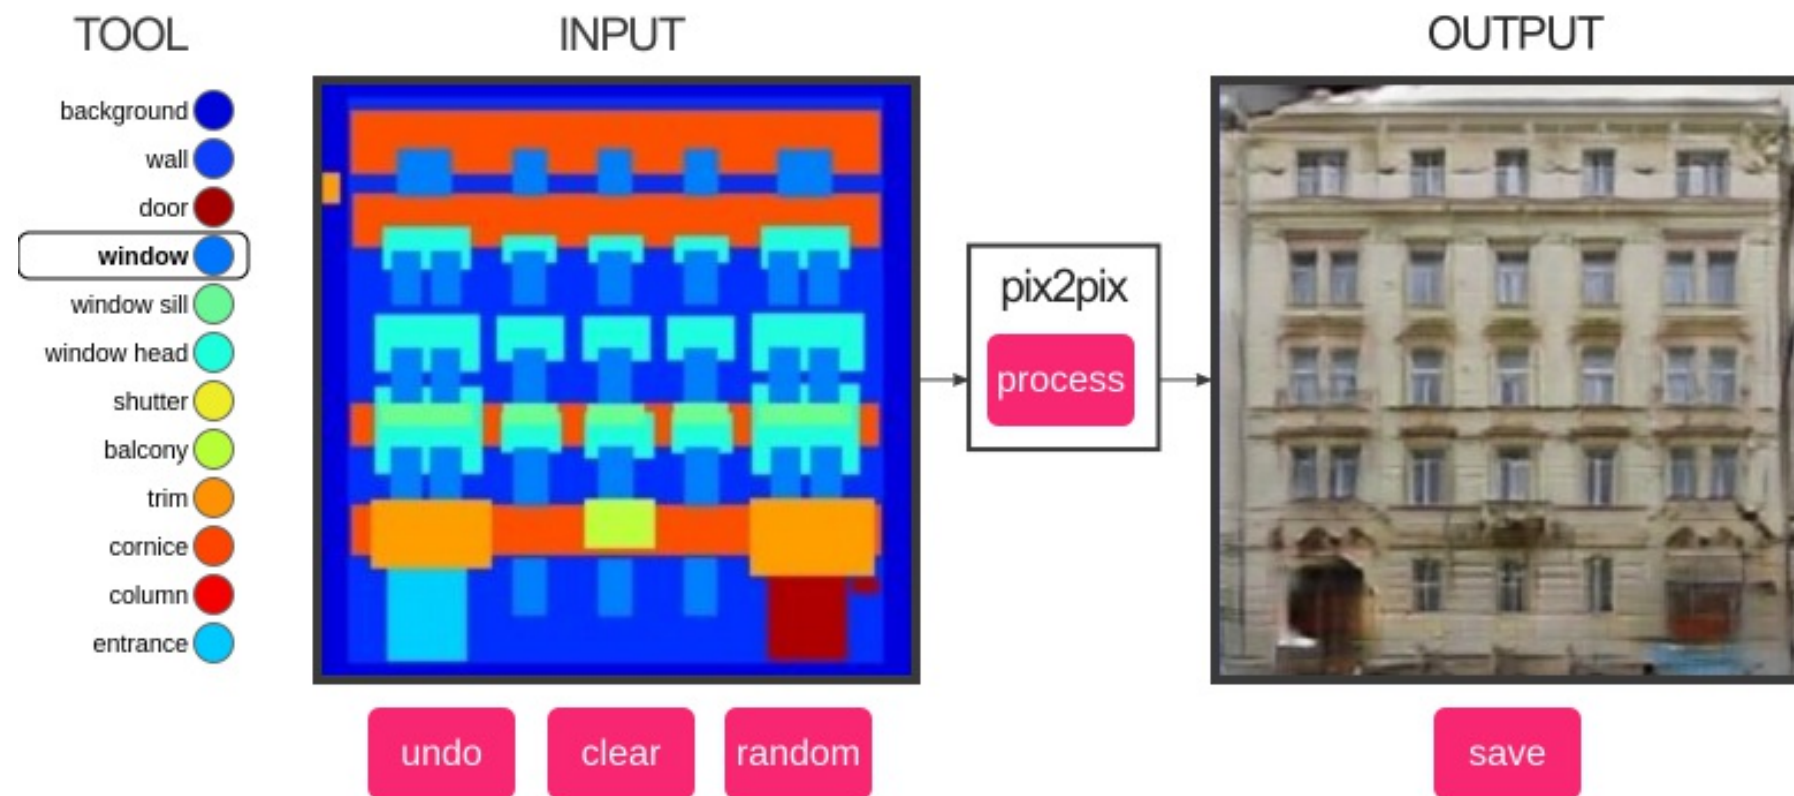
No glory in data preparation

But it must be done



Having a small dataset?

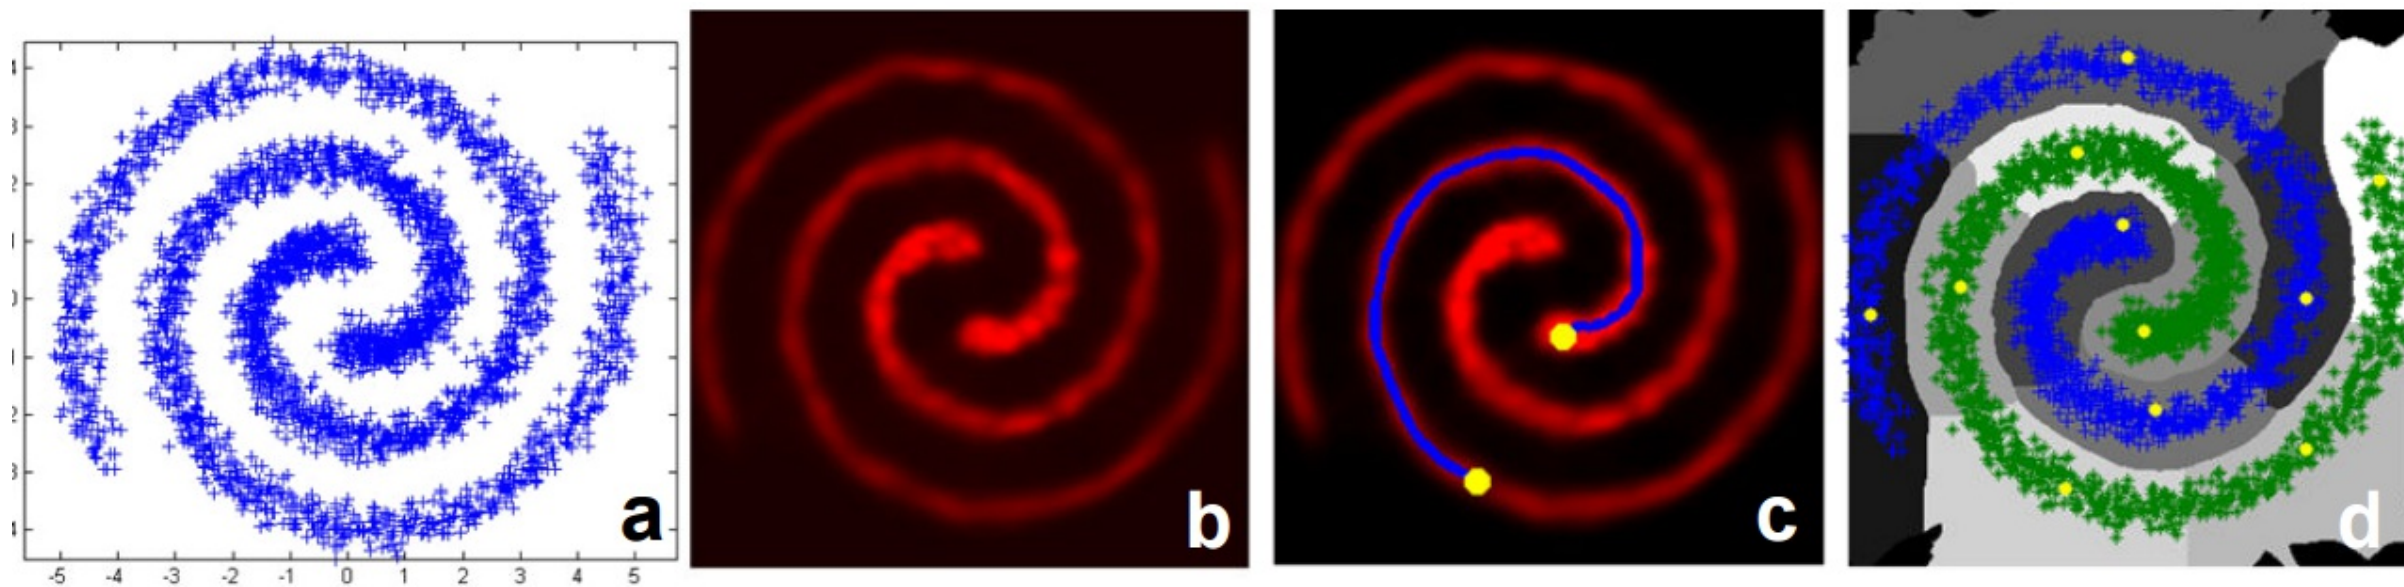
Do something about it



<https://affinelayer.com/pixsrv/>

Got unlabeled data?

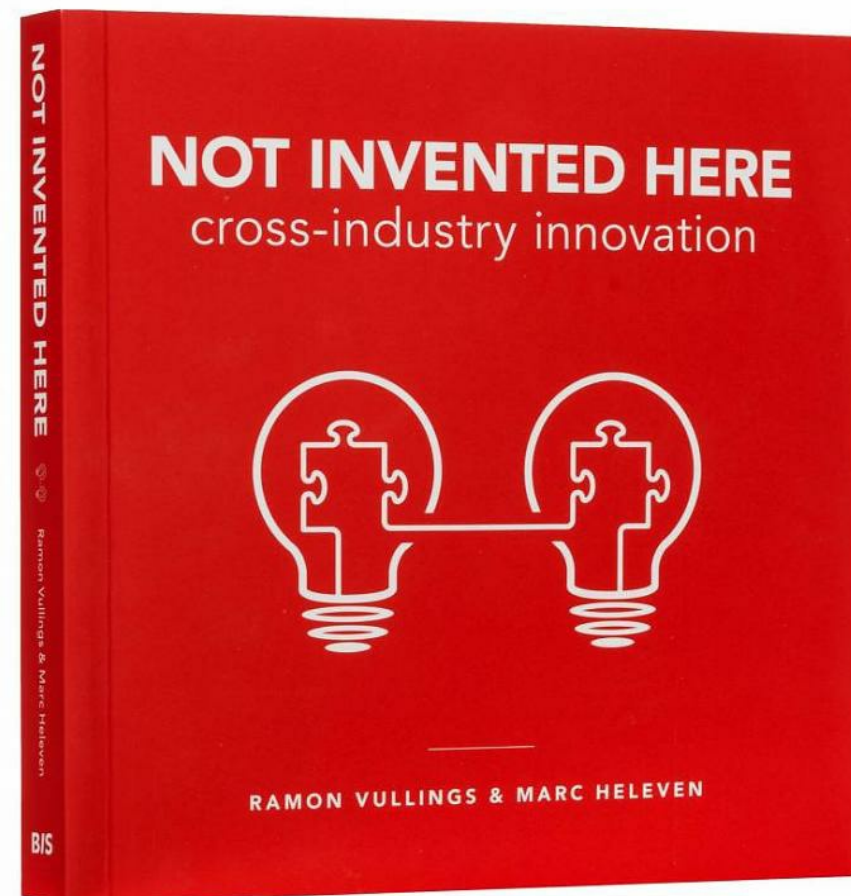
Don't be lazy, just annotate it if you can, there are tools to help you



Margeta et al. 2015, Joint work with [Inria](#) and [Microsoft Research Cambridge](#)

See **Label Propagation example** in **Scikit-learn**

Be practical
Have an open mind



Overall experience

- Remarkable ecosystem
- Fantastic community
- Build things very hard before



Takeaways

- Build something you care about
- Poke your models and learn from them
- Pick your superpower and have fun

```
# to fly  
import antigravity
```



The rod of Asclepios

Jan Margeta | jan@kardio.me | [@jmargeta](https://twitter.com/jmargeta)

KardioMe³

Thanks!

Python, numpy, MxNet, Keras, Tensorflow, scikit-learn, SimpleITK, pydicom, Flask,
Sanic, Django, PostgreSQL, ReactJS, Docker

Inria, IHU Liryc, Microsoft Research